

Lecture 4. Iterative Optimisation & Logistic Regression


COMP90051 Statistical Machine Learning

Semester 1, 2021
Lecturer: Trevor Cohn



THE UNIVERSITY OF
MELBOURNE

This lecture

- Iterative optimisation
 - * First-order method: Gradient descent
 - * Second-order: Newton-Raphson method
 -  Later: Lagrangian duality
- Logistic regression: workhorse linear classifier
 - * Possibly familiar derivation: frequentist
 - * Decision-theoretic derivation

Gradient Descent

Brief review of most basic
optimisation approach in ML

Optimisation formulations in ML

- Training = Fitting = Parameter estimation
- Typical **formulation**

$$\hat{\theta} \in \operatorname{argmin}_{\theta \in \Theta} L(\text{data}, \theta)$$

- * **argmin** because we want a **minimiser** not the **minimum**
 - Note: **argmin** can return a set (minimiser not always **unique**!)
- * Θ denotes a **model family** (including constraints)
- * L denotes some **objective function** to be optimised
 - E.g. MLE: (conditional) likelihood
 - E.g. Decision theory: (regularised) empirical risk

One we've seen: Log trick

- Instead of optimising $L(\theta)$, try convenient $\log L(\theta)$
- Why are we allowed to do this?
- **Strictly monotonic** function: $a > b \implies f(a) > f(b)$
 - * **Example**: log function!
- **Lemma**: Consider any objective function $L(\theta)$ and any strictly monotonic f . θ^* is an optimiser of $L(\theta)$ if and only if it is an optimiser of $f(L(\theta))$.
 - * Proof: Try it at home for fun!

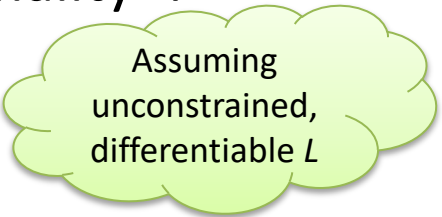
Two solution approaches

- **Analytic** (*aka* closed form) solution

- * Known only in limited number of cases

- * Use 1st-order necessary condition for optimality*:

$$\frac{\partial L}{\partial \theta_1} = \dots = \frac{\partial L}{\partial \theta_p} = 0$$



Assuming
unconstrained,
differentiable L

- **Approximate iterative** solution

1. Initialisation: choose starting guess $\theta^{(1)}$, set $i = 1$
2. Update: $\theta^{(i+1)} \leftarrow \text{SomeRule}[\theta^{(i)}]$, set $i \leftarrow i + 1$
3. Termination: decide whether to Stop
4. Go to **Step 2**
5. **Stop**: return $\hat{\theta} \approx \theta^{(i)}$

* **Note**: to check for local minimum, need positive 2nd derivative (or Hessian positive definite); this assumes unconstrained – in general need to also check boundaries. See also Lagrangian techniques later in subject.

Reminder: The gradient

- **Gradient at θ** defined as $\left[\frac{\partial L}{\partial \theta_1}, \dots, \frac{\partial L}{\partial \theta_p} \right]'$ evaluated at θ
- The gradient points to the direction of maximal change of $L(\theta)$ when departing from point θ
- Shorthand notation
 - * $\nabla L \stackrel{\text{def}}{=} \left[\frac{\partial L}{\partial \theta_1}, \dots, \frac{\partial L}{\partial \theta_p} \right]'$ computed at point θ
 - * Here ∇ is the “nabla” symbol
- **Hessian** matrix at θ : $\nabla^2 L_{ij} = \frac{\partial^2 L}{\partial \theta_i \partial \theta_j}$

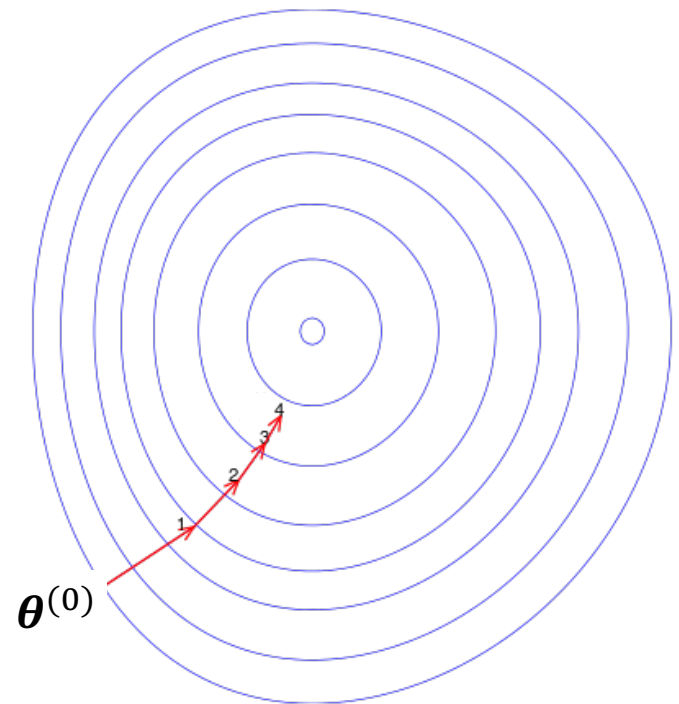


Harps, p. 984.

Gradient descent and SGD

1. Choose $\theta^{(0)}$ and some T
 2. For i from 1 to T^*
 1. $\theta^{(i)} = \theta^{(i-1)} - \eta \nabla L(\theta^{(i-1)})$
 3. Return $\hat{\theta} \approx \theta^{(T)}$
- Note: η dynamically updated per step
 - Variants: Momentum, AdaGrad, ...
 - Stochastic gradient descent: two loops
 - * Outer for loop: each loop (called **epoch**) sweeps through all training data
 - * Within each epoch, randomly shuffle training data; then for loop: do gradient steps only on **batches** of data. Batch size might be 1 or few

Assuming L is differentiable



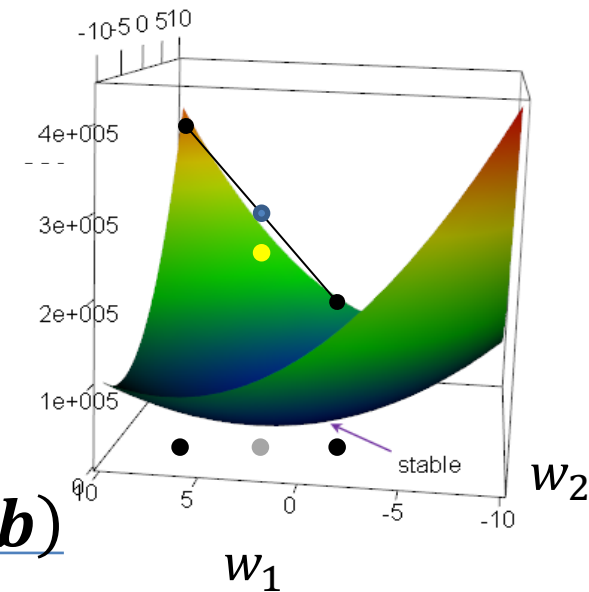
Wikimedia Commons. Authors:
Olegalexandrov, Zerodamage

*Other stopping criteria can be used

Convex objective functions

- ‘Bowl shaped’ functions
- Informally: if line segment between any two points on graph of function lies above or on graph
- Formally* $f: D \rightarrow \mathbf{R}$ is convex if $\forall \mathbf{a}, \mathbf{b} \in D, t \in [0,1]$:

$$f(t\mathbf{a} + (1-t)\mathbf{b}) \leq tf(\mathbf{a}) + (1-t)f(\mathbf{b})$$
 Strictly convex if inequality is strict ($<$)
- Gradient descent on (strictly) convex function guaranteed to find a (unique) global minimum!

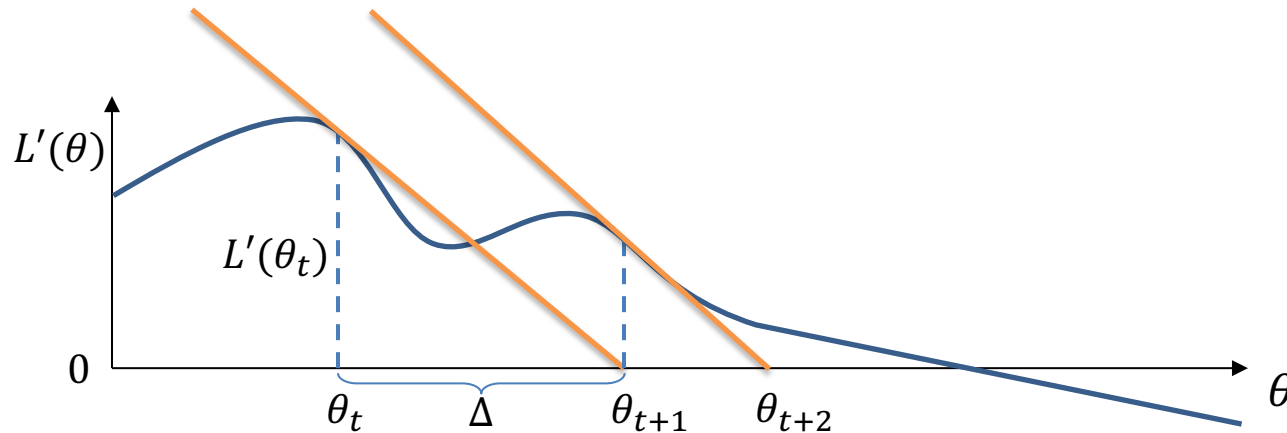


* **Aside:** Equivalently we can look to the second derivative. For f defined on scalars, it should be non-negative; for multivariate f , the Hessian matrix should be positive semi-definite (see linear algebra supplemental deck).

Newton-Raphson

A second-order method;
Successive root finding in the
objective's derivative.

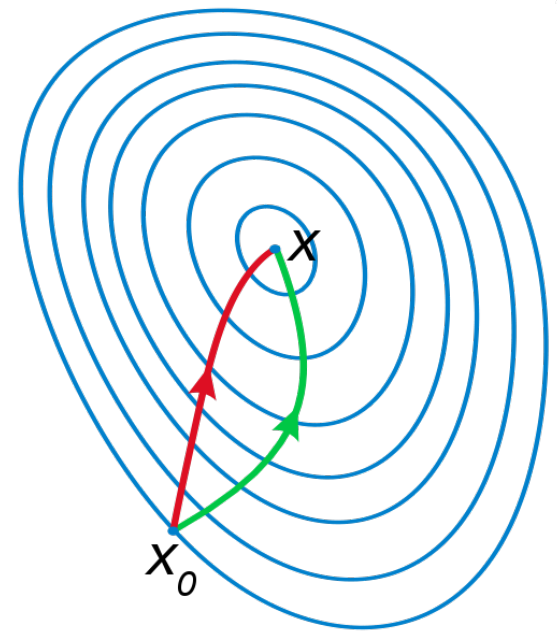
Newton-Raphson: Derivation (1D)



- Critical points of $L(\theta)$ = Zero-crossings of $L'(\theta)$
- Consider case of scalar θ . Starting at given/random θ_0 , iteratively:
 1. Fit tangent line to $L'(\theta)$ at θ_t
 2. Need to find $\theta_{t+1} = \theta_t + \Delta$ using linear approximation's zero crossing
 3. Tangent line given by derivative: rise/run = $-L''(\theta_t) = L'(\theta_t)/\Delta$
 4. Therefore iterate is $\theta_{t+1} = \theta_t - L'(\theta_t)/L''(\theta_t)$


Newton-Raphson: General case

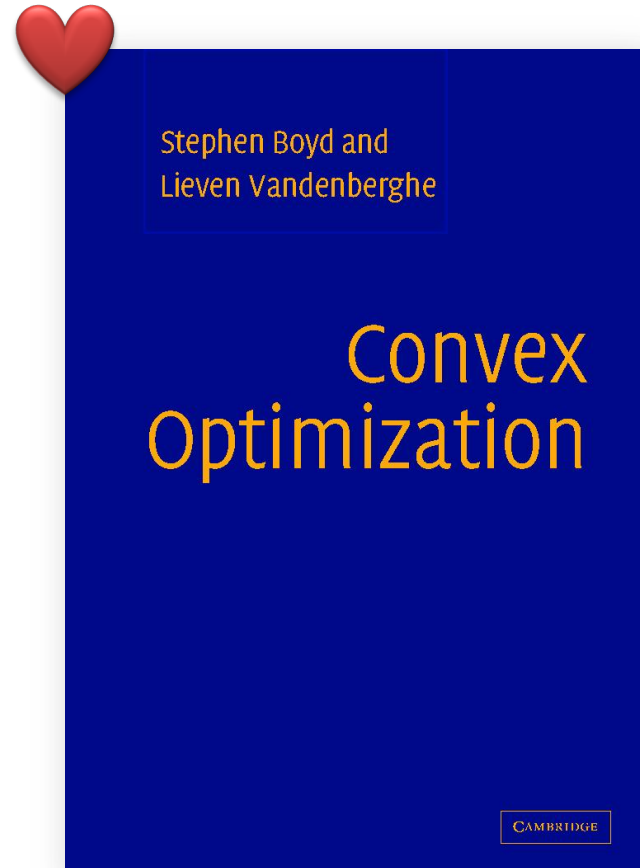
- Newton-Raphson summary
 - * Finds $L'(\theta)$ zero-crossings
 - * By successive linear approximations to $L'(\theta)$
 - * Linear approximations involve derivative of $L'(\theta)$, ie. $L''(\theta)$
- Vector-valued θ :
How to fix scalar $\theta_{t+1} = \theta_t - L'(\theta_t)/L''(\theta_t)$???
- * $L'(\theta)$ is $\nabla L(\theta)$
 - * $L''(\theta)$ is $\nabla_2 L(\theta)$
 - * Matrix division is matrix inversion
- General case: $\theta_{t+1} = \theta_t - (\nabla_2 L(\theta_t))^{-1} \nabla L(\theta_t)$
 - * Pro: May converge faster; fitting a quadratic with curvature information
 - * Con: Sometimes computationally expensive, unless approximating Hessian



public domain wikipedia

...And much much more

- What if you have constraints?
 -  See Lagrangian multipliers (let's you bring constraints into objective)
 - * Or, projected gradient descent (you iterate between GD on objective, and GD on each constraints)
- What about speed of convergence?
- Do you really need differentiable objectives? (no, subgradients)
- Are there more tricks? (Hell yeah! But outside scope here)



Free at <http://web.stanford.edu/~boyd/cvxbook/>

Mini Summary

- Iterative optimisation for ML
 - * First-order: Gradient Descent and Stochastic GD
 - * Convex objectives: Convergence to global optima
 - * Second-order: Newton-Raphson can be faster, can be expensive to build/invert full Hessian

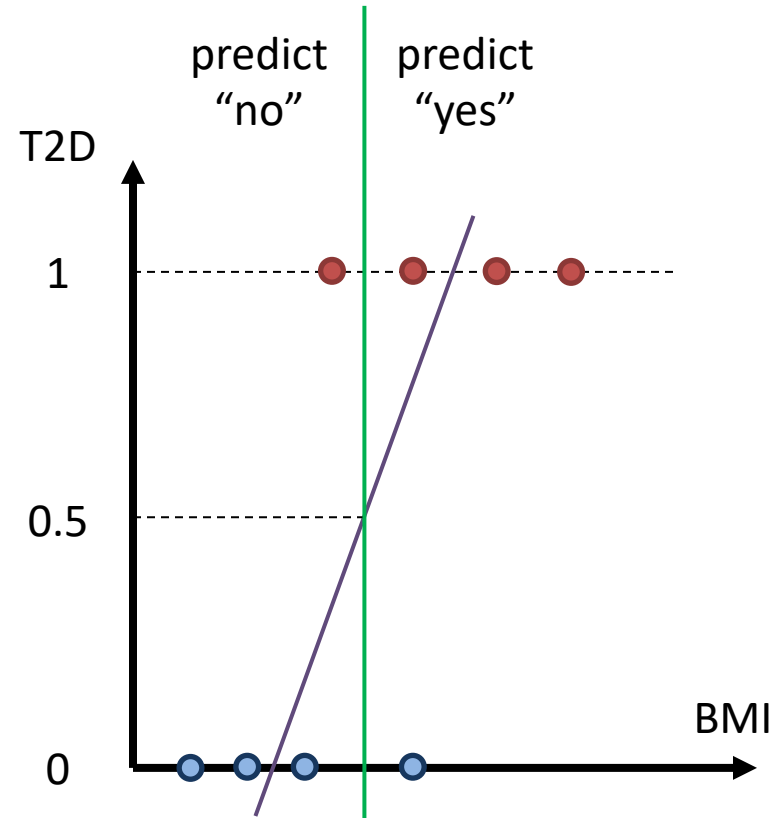
Next: Logistic regression for binary classification

Logistic Regression Model

A workhorse linear, binary classifier;
(A review for some of you; new to some.)

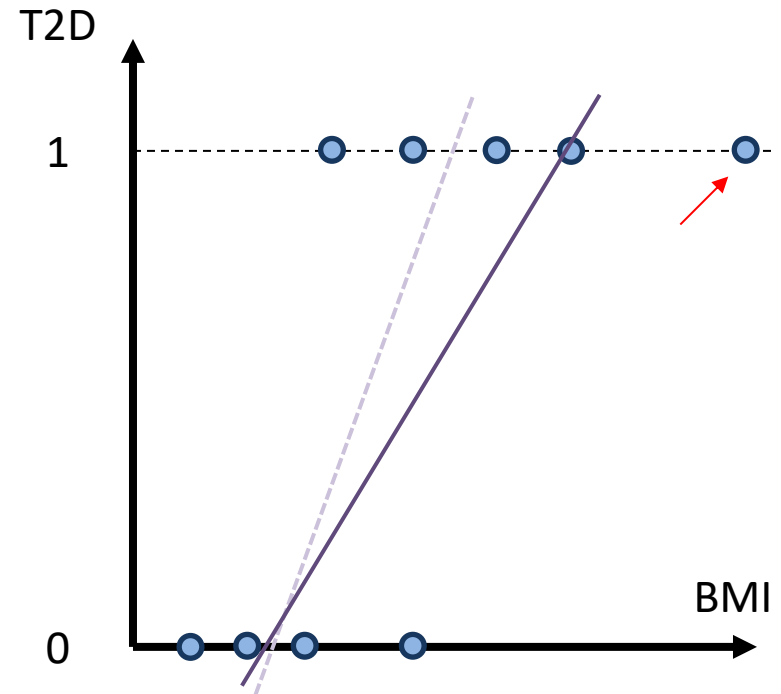
Binary classification: Example

- Example: given body mass index (BMI) does a patient have type 2 diabetes (T2D)?
- This is (supervised) **binary classification**
- One *could* use linear regression
 - * Fit a line/hyperplane to data (find weights \mathbf{w})
 - * Denote $s \equiv \mathbf{x}'\mathbf{w}$
 - * Predict “Yes” if $s \geq 0.5$
 - * Predict “No” if $s < 0.5$



Why not linear regression

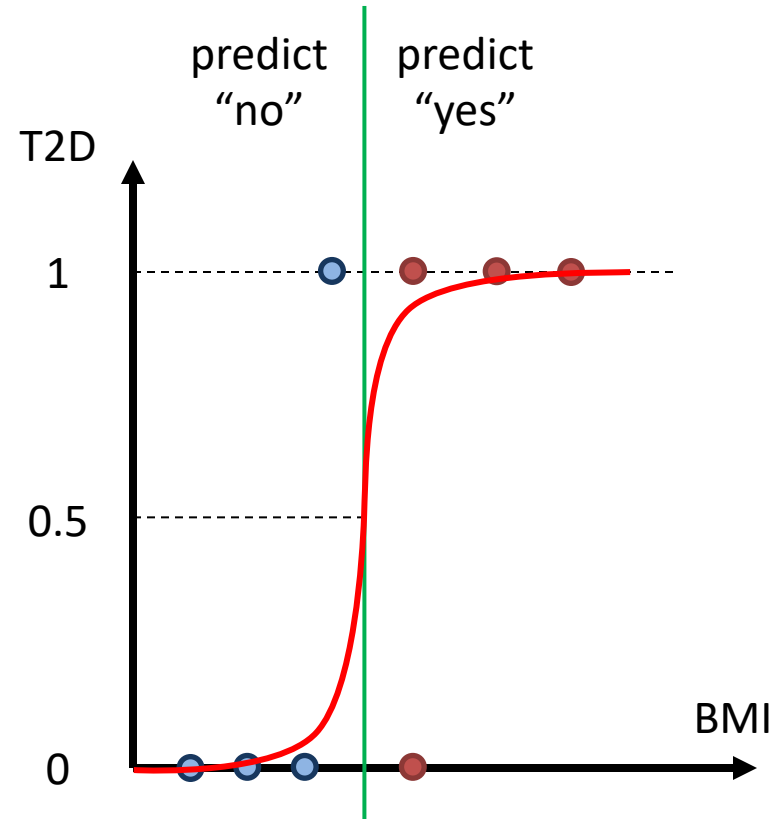
- Due to the square loss, points far from boundary have loss squared – even if they’re confidently correct!
- Such “outliers” will “pull at” the linear regression
- Overall, the least-squares criterion looks unnatural in this setting



Logistic regression model

- Probabilistic approach to classification
 - * $P(Y = 1|\mathbf{x}) = f(\mathbf{x}) = ?$
 - * Use a linear function? E.g., $s(\mathbf{x}) = \mathbf{x}'\mathbf{w}$
- Problem: the probability needs to be between 0 and 1.
- **Logistic** function $f(s) = \frac{1}{1+\exp(-s)}$
- **Logistic regression model**

$$P(Y = 1|\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{x}'\mathbf{w})}$$



How is logistic regression *linear*?

- Logistic regression model:

$$P(Y = 1|\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{x}'\mathbf{w})}$$

- Classification rule:

if $(P(Y = 1|\mathbf{x}) > \frac{1}{2})$ then class “1”, else class “0”

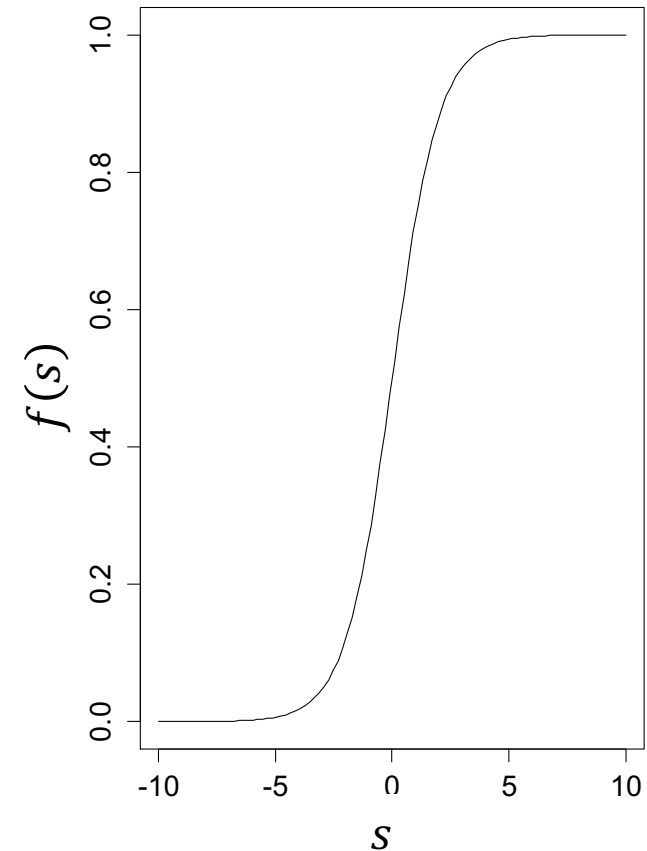
- Decision boundary is the set of \mathbf{x} 's such that:

$$\frac{1}{1 + \exp(-\mathbf{x}'\mathbf{w})} = \frac{1}{2}$$

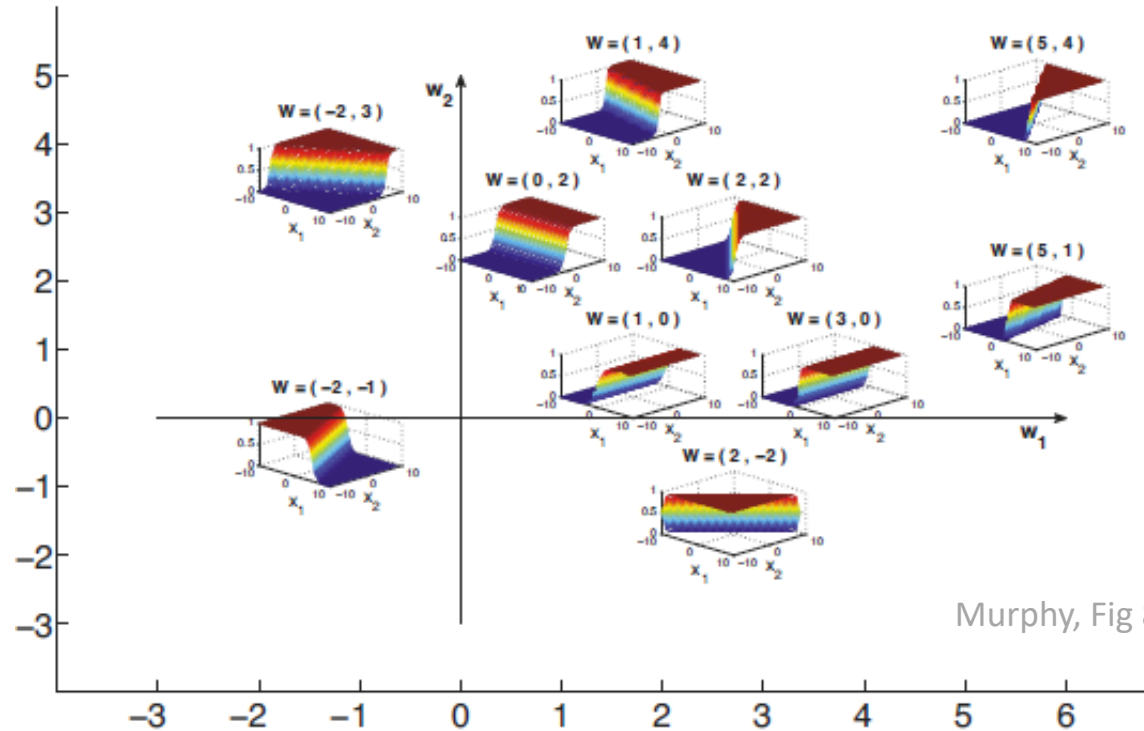
$$\exp(-\mathbf{w}'\mathbf{x}) = 1$$

$$\mathbf{w}'\mathbf{x} = 0$$

Logistic function



Effect of parameter vector (2D problem)



Murphy, Fig 8.1, p246

- **Decision boundary** is the line where $P(Y = 1|\mathbf{x}) = 0.5$
 - * In higher dimensional problems, the decision boundary is a plane or hyperplane
- Vector \mathbf{w} is perpendicular to the decision boundary (see linear algebra review topic)
 - * That is, \mathbf{w} is a normal to the decision boundary
 - * Note: in this illustration we assume $w_0 = 0$ for simplicity

Linear vs. logistic probabilistic models

- **Linear regression** assumes a Normal distribution with a fixed variance and mean given by linear model

$$p(y|\mathbf{x}) = \text{Normal}(\mathbf{x}'\mathbf{w}, \sigma^2)$$

- **Logistic regression** assumes a Bernoulli distribution with parameter given by logistic transform of linear model

$$p(y|\mathbf{x}) = \text{Bernoulli}(\text{logistic}(\mathbf{x}'\mathbf{w}))$$

- Recall that **Bernoulli distribution** is defined as

$$p(1) = \theta \text{ and } p(0) = 1 - \theta \text{ for } \theta \in [0,1]$$

- Equivalently $p(y) = \theta^y (1 - \theta)^{(1-y)}$ for $y \in \{0,1\}$

Training as Max-Likelihood Estimation

- Assuming independence, probability of data

$$p(y_1, \dots, y_n | \mathbf{x}_1, \dots, \mathbf{x}_n) = \prod_{i=1}^n p(y_i | \mathbf{x}_i)$$

- Assuming Bernoulli distribution we have

$$p(y_i | \mathbf{x}_i) = (\theta(\mathbf{x}_i))^{y_i} (1 - \theta(\mathbf{x}_i))^{1-y_i}$$

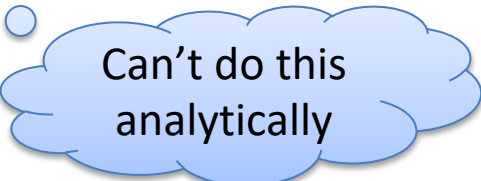
$$\text{where } \theta(\mathbf{x}_i) = \frac{1}{1 + \exp(-\mathbf{x}_i' \mathbf{w})}$$

- Training: maximise this expression wrt weights \mathbf{w}

Apply log trick, simplify

- Instead of maximising likelihood, maximise its logarithm

$$\begin{aligned}\log \left(\prod_{i=1}^n p(y_i | \mathbf{x}_i) \right) &= \sum_{i=1}^n \log p(y_i | \mathbf{x}_i) \\ &= \sum_{i=1}^n \log \left((\theta(\mathbf{x}_i))^{y_i} (1 - \theta(\mathbf{x}_i))^{1-y_i} \right) \\ &= \sum_{i=1}^n (y_i \log(\theta(\mathbf{x}_i)) + (1 - y_i) \log(1 - \theta(\mathbf{x}_i))) \\ &= \sum_{i=1}^n ((y_i - 1) \mathbf{x}_i' \mathbf{w} - \log(1 + \exp(-\mathbf{x}_i' \mathbf{w})))\end{aligned}$$



Can't do this
analytically

Training as Iterative Optimisation

- Training logistic regression: \mathbf{w} maximising log-likelihood $L(\mathbf{w})$ or cross-entropy loss
- **Bad news:** No closed form solution
- **Good news:** Problem is strictly convex, if no irrelevant features \rightarrow convergence!

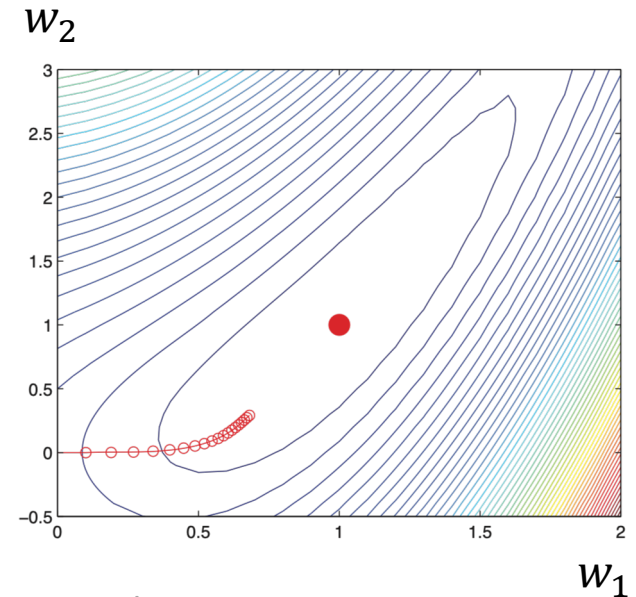


Look ahead: regularisation for irrelevant features

How does gradient descent work?

- simply take gradient of log-likelihood, i.e.,

$$\nabla L(\mathbf{w}) = \sum_{i=1}^n (y_i - \theta(\mathbf{x}_i)) \mathbf{x}_i$$
- plug into favourite iterative optimiser
 GD/SGD/Adagrad/Adam/BFGS/...



Murphy, Fig 8.3, p247

Logistic Regression: Decision-Theoretic View

Via cross-entropy loss

Background: Cross entropy

- Cross entropy is an information-theoretic method for comparing two distributions
- Cross entropy is a measure of a **divergence** between reference distribution $g_{ref}(a)$ and estimated distribution $g_{est}(a)$. For discrete distributions:

$$H(g_{ref}, g_{est}) = - \sum_{a \in A} g_{ref}(a) \log g_{est}(a)$$

A is support of the distributions, e.g., $A = \{0,1\}$

Training as cross-entropy minimisation

- Consider log-likelihood for a single data point
$$\log p(y_i | \mathbf{x}_i) = y_i \log(\theta(\mathbf{x}_i)) + (1 - y_i) \log(1 - \theta(\mathbf{x}_i))$$
- Cross entropy $H(g_{ref}, g_{est}) = -\sum_a g_{ref}(a) \log g_{est}(a)$
 - * If reference (true) distribution is

$$g_{ref}(1) = y_i \text{ and } g_{ref}(0) = 1 - y_i$$

- * With logistic regression estimating this distribution as

$$g_{est}(1) = \theta(\mathbf{x}_i) \text{ and } g_{est}(0) = 1 - \theta(\mathbf{x}_i)$$

It finds \mathbf{w} that minimises sum of cross entropies per training point

Mini Summary

- Logistic regression *formulation*
 - * A workhorse linear binary classifier (but can use basis function to become non-linear)
 - * Frequentist: Bernoulli label with coin bias logistic-linear in \mathbf{x}
 - * Decision theory: Minimising cross entropy with labels

Next time: Regularised linear regression for avoiding overfitting and ill-posed optimisation