

# Lecture 21. HMMs and Message Passing

COMP90051 Statistical Machine Learning

Semester 1, 2021  
Lecturer: Trevor Cohn



# This lecture


- Hidden Markov models – detailed PGM case study
  - \* Brief recap of model
  - \* “Evaluation”: Forward-Background Algorithm = elimination
  - \* “Learning”: Baum Welch = MLE
  - \* “Decoding”: Viterbi = elimination variant with  $\text{sum} \rightarrow \text{max}$
- Message passing
  - \* Sum-product generalises elimination algorithm
  - \* Variants for ring operators, max-product for Viterbi
  - \* Factor graphs

# Hidden Markov Models

*Model of choice for sequential data. A form of clustering for discrete time series.*

# HMM Formulation

- Formulated as directed PGM
  - therefore joint expressed as



$$P(\mathbf{o}, \mathbf{q}) = P(q_1)P(o_1|q_1) \prod_{i=2}^T P(q_i|q_{i-1})P(o_i|q_i)$$

- bold** variables are shorthand for vector of  $T$  values
- Parameters (for *homogenous* HMM)

$A = \{a_{ij}\}$	transition probability matrix; $\forall i : \sum_j a_{ij} = 1$
$B = \{b_i(o_k)\}$	output probability matrix; $\forall i : \sum_k b_i(o_k) = 1$
$\Pi = \{\pi_i\}$	the initial state distribution; $\sum_i \pi_i = 1$

# Fundamental HMM Tasks

HMM Task	PGM Task
<b>Evaluation.</b> Given an HMM $\mu$ and observation sequence $\mathbf{o}$ , determine likelihood $\Pr(\mathbf{o} \mu)$	Probabilistic inference
<b>Decoding.</b> Given an HMM $\mu$ and observation sequence $\mathbf{o}$ , determine most probable hidden state sequence $\mathbf{q}$	MAP point estimate
<b>Learning.</b> Given an observation sequence $\mathbf{o}$ and set of states, learn parameters $A, B, \Pi$	Statistical inference

# “Evaluation” a.k.a. marginalisation

- Compute prob. of observations  $\mathbf{o}$  by summing out  $\mathbf{q}$

$$\begin{aligned} P(\mathbf{o}|\mu) &= \sum_{\mathbf{q}} P(\mathbf{o}, \mathbf{q}|\mu) \\ &= \sum_{q_1} \sum_{q_2} \dots \sum_{q_T} P(q_1)P(o_1|q_1)P(q_2|q_1)P(o_2|q_2) \dots P(q_T|q_{T-1})P(o_T|q_T) \end{aligned}$$

- Make this more efficient by moving the sums

$$P(\mathbf{o}|\mu) = \sum_{q_1} P(q_1)P(o_1|q_1) \sum_{q_2} P(q_2|q_1)P(o_2|q_2) \dots \sum_{q_T} P(q_T|q_{T-1})P(o_T|q_T)$$

- Déjà vu? Maybe we could do var. elimination...

# Elimination = Backward Algorithm

$$P(\mathbf{o}|\mu) = \sum_{q_1} P(q_1)P(o_1|q_1) \sum_{q_2} P(q_2|q_1)P(o_2|q_2) \dots \sum_{q_T} P(q_T|q_{T-1})P(o_T|q_T)$$

Eliminate  $q_T$

$$m_{T \rightarrow T-1}(q_{T-1})$$


...

Eliminate  $q_2$

$$m_{2 \rightarrow 1}(q_1)$$

“Eliminate”  $q_1$

$$P(\mathbf{o}|\mu) = \sum_{q_1} P(q_1)P(o_1|q_1)m_{2 \rightarrow 1}(q_1)$$



# Elimination = Forward Algorithm

$$P(\mathbf{o}|\mu) = \sum_{q_T} P(o_T|q_T) \sum_{q_{T-1}} P(q_T|q_{T-1}) P(o_T|q_T) \dots \sum_{q_1} P(q_2|q_1) P(q_1) P(o_1|q_1)$$

Eliminate  $q_1$

...


Eliminate  $q_{T-1}$

“Eliminate”  $q_T$

$m_{1 \rightarrow 2}(q_2)$

$m_{T-1 \rightarrow T}(q_T)$


$$P(\mathbf{o}|\mu) = \sum_{q_1} P(o_T|q_T) m_{T-1 \rightarrow T}(q_T)$$





# Variable elimination perspective

- Both algorithms are just *variable elimination* using different orderings
  - \*  $q_T \dots q_1 \rightarrow$  backward algorithm
  - \*  $q_1 \dots q_T \rightarrow$  forward algorithm
  - \* both have time complexity  $O(TL^2)$  for  $L$  the label set size
- Can use either to compute  $P(\mathbf{o})$
- Even though these are just instances of elimination, they pre-date general PGM inference.
  - \* E.g. called the “forward-backward algorithm”
  - \* Both directions useful in statistical inference (next)



# Quick aside: Viterbi

- What happens if we switch sums for max?
  - \* Viola! Finds  $\max_{\mathbf{q}} P(\mathbf{q}|\mathbf{o})$ ; and with some book-keeping can recover argmax
  - \* More on this later...

COMP90051 Statistical Machine Learning

### Elimination = Forward Algorithm

~~$P(\mathbf{o}|\mathbf{q}) = \sum_{q_T} P(o_T|q_T) \sum_{q_{T-1}} P(q_T|q_{T-1}) P(o_T|q_T) \dots \sum_{q_1} P(q_2|q_1) P(q_1) P(o_1|q_1)$~~

Eliminate  $q_1$   
 ...  
 Eliminate  $q_{T-1}$   
 "Eliminate"  $q_T$

~~$P(\mathbf{o}|\mathbf{q}) = \sum_{q_1} P(o_T|q_T) m_{T-1 \rightarrow T}(q_T)$~~

$m_{1 \rightarrow 2}(q_2)$   
 $m_{T-1 \rightarrow T}(q_T)$

$q_1 \rightarrow q_2 \rightarrow \dots \rightarrow q_T$

8

# Mini Summary

- HMM
  - \* Powerful and versatile model
  - \* “Algorithms” for HMM just instances of PGM machinery
- Evaluation by Forward / Backward
  - \* Just elimination by two different orderings

**Next time:** Statistical inference (learning) example of EM

# Statistical Inference (Learning)

- Learn parameters  $\mu = (A, B, \pi)$ , given observation sequence  $\mathbf{o}$
- Called “**Baum Welch**” algorithm which uses **EM**<sup>\*</sup> to approximate MLE,  $\text{argmax}_{\mu} P(\mathbf{o} | \mu)$ :
  1. initialise  $\mu^1$ , let  $j=1$
  2. compute expected marginal distributions  $P(q_t | \mathbf{o}, \mu^j)$  for all  $t$ ; and  $P(q_{t-1}, q_t | \mathbf{o}, \mu^j)$  for  $t=2..T$
  3. fit model  $\mu^{j+1}$  based on expectations
  4. repeat from step 2, with  $j=j+1$
- Expectations (2) computed using **forward-backward**

E step  
M step

\* Expectation-Maximisation (EM) is coming up

# Forward-Backward for $P(q_i|\mathbf{o})$

- Forward-Backward gives: messages,  $P(\mathbf{o})$

- Bayes rule:  $P(q_i|\mathbf{o}) = \frac{P(q_i, \mathbf{o})}{P(\mathbf{o})}$

- Marginalisation:  $P(q_i, \mathbf{o}) = \sum_{q_1, \dots, q_{i-1}, q_{i+1}, \dots, q_T} P(\mathbf{q}, \mathbf{o})$

$$\begin{aligned}
 &= \left( \sum_{q_1, \dots, q_{i-1}} P(o_1, \dots, o_{i-1}, q_1, \dots, q_i) \right) P(o_i | q_i) \left( \sum_{q_{i+1}, \dots, q_T} P(o_{i+1}, \dots, o_T, q_{i+1}, \dots, q_T | q_i) \right) \\
 &= m_{i-1 \rightarrow i}(q_i) P(o_i | q_i) m_{i+1 \rightarrow i}(q_i)
 \end{aligned}$$

$$P(q_i | \mathbf{o}) = \frac{1}{P(\mathbf{o})} \underbrace{m_{i-1 \rightarrow i}(q_i)}_{\text{forward}} P(o_i | q_i) \underbrace{m_{i+1 \rightarrow i}(q_i)}_{\text{backward}}$$

# Forward-Backward for $P(q_{i-1}, q_i | \mathbf{o})$

- Similar pattern:  $P(q_{i-1}, q_i | \mathbf{o}) = \frac{P(q_{i-1}, q_i, \mathbf{o})}{P(\mathbf{o})}$
- Marginalisation:  $P(q_{i-1}, q_i, \mathbf{o}) = \sum_{q_1, \dots, q_{i-2}, q_{i+1}, \dots, q_T} P(\mathbf{q}, \mathbf{o})$   

$$= \left( \sum_{q_1, \dots, q_{i-2}} P(o_1, \dots, o_{i-2}, q_1, \dots, q_{i-1}) \right) P(o_{i-1} | q_{i-1}) P(q_i | q_{i-1}) P(o_i | q_i) \left( \sum_{q_{i+1}, \dots, q_T} P(o_{i+1}, \dots, o_T, q_{i+1}, \dots, q_T | q_i) \right)$$

$$= m_{i-2 \rightarrow i-1}(q_{i-1}) P(o_{i-1} | q_{i-1}) P(q_i | q_{i-1}) P(o_i | q_i) m_{i+1 \rightarrow i}(q_i)$$

$$\frac{1}{P(\mathbf{o})} m_{i-2 \rightarrow i-1}(q_{i-1}) P(o_{i-1} | q_{i-1}) P(q_i | q_{i-1}) P(o_i | q_i) m_{i+1 \rightarrow i}(q_i)$$

forward
backward

# Mini Summary

- Statistical inference for HMMs
  - \* “Just” learning or MLE as we’re frequentist here
  - \* Unobserved random variables means: EM (more later on)
  - \* Maximisation step: looks like MLE – nothing new
  - \* Expectation step: achieved by forward-backward messages
- “Baum-Welch” is the original name of this algorithm

**Next time:** Message passing a little more generally


# Message Passing

*Sum-product algorithm for efficiently computing marginal distributions over trees. An extension of variable elimination algorithm.*




# Inference as message passing

- Each  $m$  can be considered as a **message** which summarises the effect of the rest of the graph on the current node marginal.
  - \* *Inference = passing messages between all nodes*



# Inference as message passing


- Messages vector valued, i.e., function of target label
- Messages defined recursively: left to right, or right to left for the HMM



# Sum-product algorithm

- Message passing in more general graphs

- \* applies to chains, trees and poly-trees (D-PGMs with  $>1$  parent)
- \* 'sum-product' derives from:
  - **product** = product of incoming messages
  - **sum** = summing out the effect of rv(s) *aka* elimination




- Algorithm supports other operations (semi-rings\*)

- \* e.g., max-product, swapping **sum** for **max**
- \* **Viterbi algorithm** is the max-product variant of forward algorithm, solves the  $\operatorname{argmax}_{\mathbf{q}} P(\mathbf{q}|\mathbf{o})$




\* A ring is an algebraic structure generalizing addition/multiplication on reals. Semi-ring relaxes requirement of additive inverse.

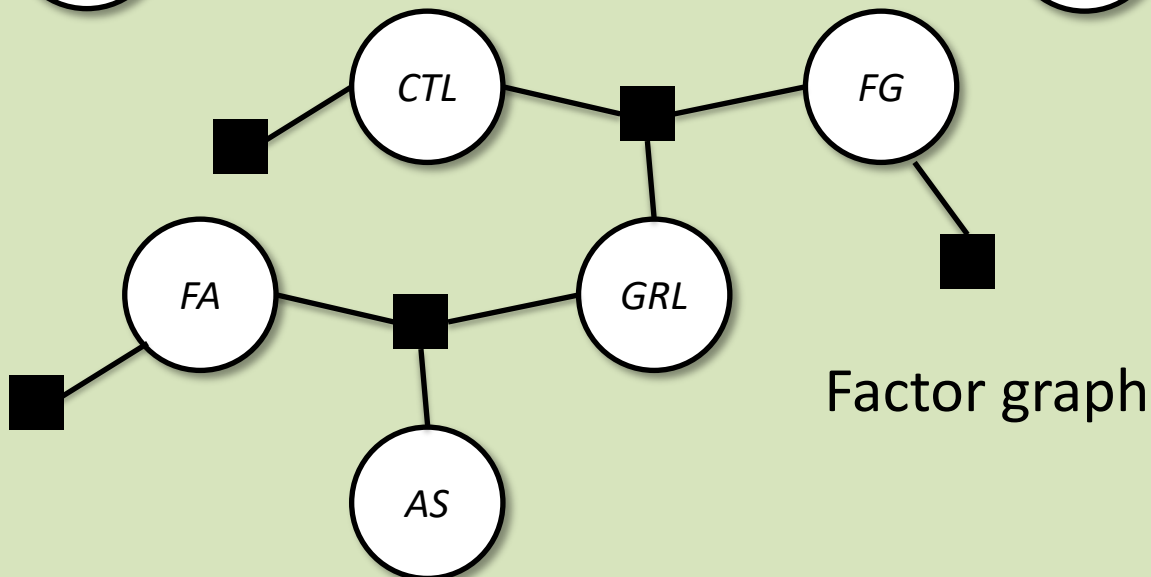
# Application to Directed PGMS



Directed PGM



Undirected  
"moralised"  
PGM

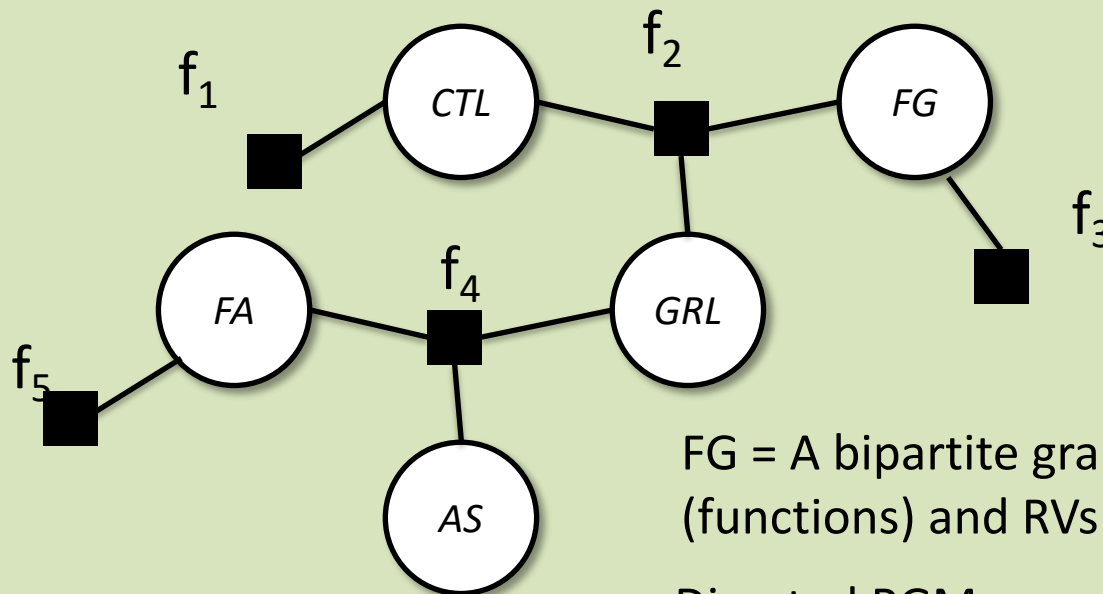


Factor graph

# Factor graphs

$$f_1(CTL) = P(CTL)$$

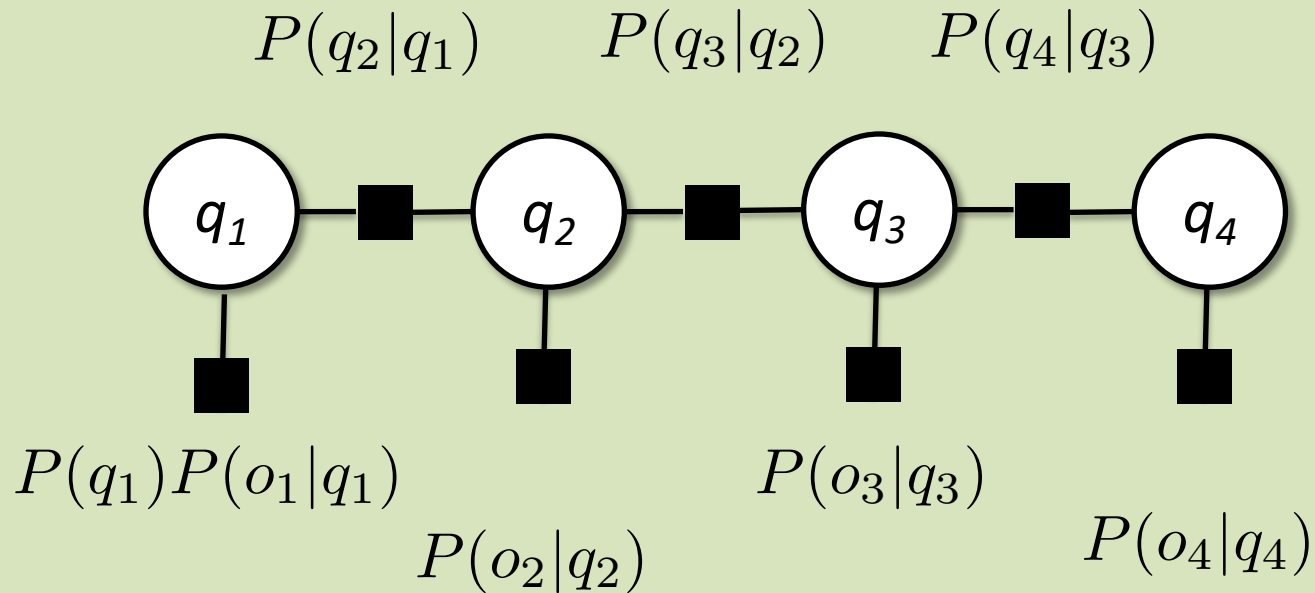
$$f_2(CTL, GRL, FG) = P(GRL|CTL, FG)$$



FG = A bipartite graph, with factors (functions) and RVs

Directed PGMs result in tree-structured FG


# Factor graph for the HMM



Effect of observed nodes incorporated into unary factors

# Advantage of Factor Graphs

- Factorisation is a central idea
- D-PGMs and U-PGMs not able to fully represent arbitrary factorisations of joints



$$p(x, y, z) \propto \varphi(x, y)\varphi(y, z)\varphi(z, x)$$
$$p(x, y, z) \propto \varphi(x, y, z)$$

- Better representation of factorisations has advantages; factor graphs are general.

# Sum-Product over Factor Graphs

- Two types of messages :
  - \* between factors and RVs; and between RVs and factors
  - \* they summarise a complete sub-graph

- E.g.,

$$m_{f_2 \rightarrow GRL}(GRL) = \sum_{CTL} \sum_{FG} f_2(GRL, CTL, FG) m_{CTL \rightarrow f_2}(CTL) m_{FG \rightarrow f_2}(FG)$$

- Structure inference as “gather-and-distribute”
  - \* gather messages from leaves of tree towards root
  - \* then propagate message back down from root to leaves



# Summary

- HMMs as example PGMs
  - \* formulation as PGM
  - \* independence assumptions
  - \* probabilistic inference using forward-backward
  - \* statistical inference using expectation-maximization
  - \* decoding as max-product
- Message passing: general inference method for U-PGMs
  - \* sum-product & max-product
  - \* factor graphs

**Next time:** Gaussian mixture models and EM