



Workshop 6

COMP90051 Statistical Machine Learning

Semester 1, 2021


Learning Outcomes

At the end of this workshop you should:

1. Be familiar with **Convolutional Neural Network** and be able to implement it.
2. Be able to implement autoencoder

Image classification task

- CIFAR 10 - Image classification
- 10 labels
 - * Multi-class
 - * classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')



How is data represented in neural nets?

- Primary data structure is the **tensor**—a fancy name for a multi-dimensional array
- Can be used to represent trainable weights, hyperparameters, data flowing through the network, etc.
- E.g. an **image** can be represented as a **3d-tensor**: 2 dimensions for horizontal/vertical pixels + 1 dimension for the RGB channels






Image Tensor


2d tensor (e.g. grayscale image)



3d tensor (e.g. rgb image)



4d tensor (e.g. sequence of rgb images)



Feature for logistic regression



```
tensor([[0.3647, 0.3608, 0.3647, ..., 0.3804, 0.3882, 0.3843],
        [0.3725, 0.3608, 0.3608, ..., 0.4000, 0.3882, 0.3529],
        [0.4000, 0.3843, 0.3765, ..., 0.4275, 0.4039, 0.3608],
        ...,
        [0.7961, 0.7882, 0.8000, ..., 0.8157, 0.8314, 0.8235],
        [0.8000, 0.7922, 0.7882, ..., 0.8431, 0.8314, 0.7804],
        [0.7804, 0.7765, 0.7765, ..., 0.7961, 0.7725, 0.7451]])
```

32*32 matrix



```
tensor([[0.4431, 0.4392, 0.4392, ..., 0.4784, 0.4627, 0.4431],
        [0.4588, 0.4588, 0.4627, ..., 0.4980, 0.4824, 0.4392],
        [0.4784, 0.4863, 0.4863, ..., 0.5137, 0.5020, 0.4549],
        ...,
        [0.7725, 0.7647, 0.7765, ..., 0.7882, 0.8039, 0.7922],
        [0.7765, 0.7686, 0.7647, ..., 0.8157, 0.8039, 0.7529],
        [0.7569, 0.7529, 0.7529, ..., 0.7686, 0.7451, 0.7176]])
```

32*32 matrix



```
tensor([[0.7176, 0.7098, 0.7059, ..., 0.7294, 0.7255, 0.7020],
        [0.7412, 0.7098, 0.7059, ..., 0.7608, 0.7255, 0.6902],
        [0.6980, 0.6863, 0.7216, ..., 0.7176, 0.6824, 0.6745],
        ...,
        [0.7059, 0.7020, 0.7098, ..., 0.7255, 0.7412, 0.7294],
        [0.7137, 0.7059, 0.7020, ..., 0.7529, 0.7412, 0.6902],
        [0.6941, 0.6902, 0.6902, ..., 0.7059, 0.6824, 0.6549]])
```

32*32 matrix



flatten

```
tensor([[[[0.3647, 0.3608, 0.3647, ..., 0.7059, 0.6824, 0.6549]]]])
```

3*32*32
tensor

tensor([[[0.3647, 0.3608, 0.3647, ..., 0.7059, 0.6824, 0.6549]]])

1*3072 vector

Feature for CNN



```
tensor([[[0.3647, 0.3608, 0.3647, ..., 0.3804, 0.3882, 0.3843],
         [0.3725, 0.3608, 0.3608, ..., 0.4000, 0.3882, 0.3529],
         [0.4000, 0.3843, 0.3765, ..., 0.4275, 0.4039, 0.3608],
         ...,
         [0.7961, 0.7882, 0.8000, ..., 0.8157, 0.8314, 0.8235],
         [0.8000, 0.7922, 0.7882, ..., 0.8431, 0.8314, 0.7804],
         [0.7804, 0.7765, 0.7765, ..., 0.7961, 0.7725, 0.7451]])])
```

32*32 matrix




```
tensor([[[0.4431, 0.4392, 0.4392, ..., 0.4784, 0.4627, 0.4431],
         [0.4588, 0.4588, 0.4627, ..., 0.4980, 0.4824, 0.4392],
         [0.4784, 0.4863, 0.4863, ..., 0.5137, 0.5020, 0.4549],
         ...,
         [0.7725, 0.7647, 0.7765, ..., 0.7882, 0.8039, 0.7922],
         [0.7765, 0.7686, 0.7647, ..., 0.8157, 0.8039, 0.7529],
         [0.7569, 0.7529, 0.7529, ..., 0.7686, 0.7451, 0.7176]])])
```

32*32 matrix



```
tensor([[[0.7176, 0.7098, 0.7059, ..., 0.7294, 0.7255, 0.7020],
         [0.7412, 0.7098, 0.7059, ..., 0.7608, 0.7255, 0.6902],
         [0.6980, 0.6863, 0.7216, ..., 0.7176, 0.6824, 0.6745],
         ...,
         [0.7059, 0.7020, 0.7098, ..., 0.7255, 0.7412, 0.7294],
         [0.7137, 0.7059, 0.7020, ..., 0.7529, 0.7412, 0.6902],
         [0.6941, 0.6902, 0.6902, ..., 0.7059, 0.6824, 0.6549]])])
```

32*32 matrix

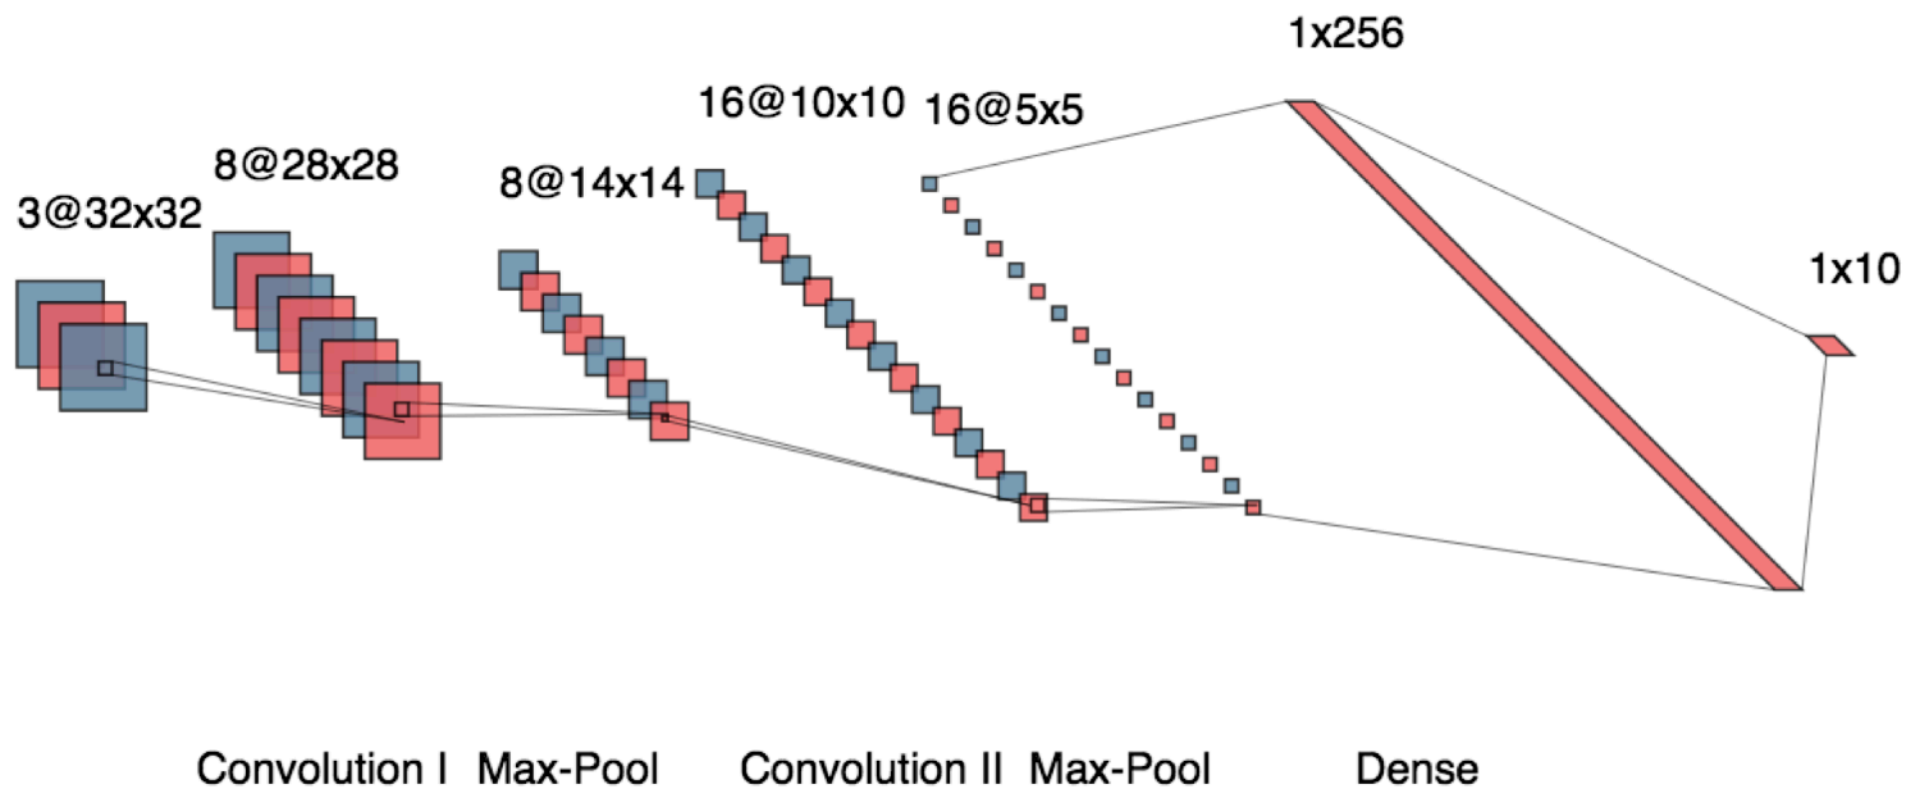


3*32*32
tensor



Then?


Architecture of our model



Components of a CNN


- Convolutional layers
 - * Complex input representations based on convolution operation
 - * Filter weights are learned from training data
- Downsampling, usually via Max Pooling
 - * Re-scales to smaller resolution, limits parameter explosion
- Fully connected parts and output layer
 - * Merges representations together

Number of Parameters




- ANN
- How many weights?

Number of Parameters




- ANN
- Consider bias
- How many weights?

Number of Parameters



- CNN
- 1D Kernel size = 3
- How many weights?

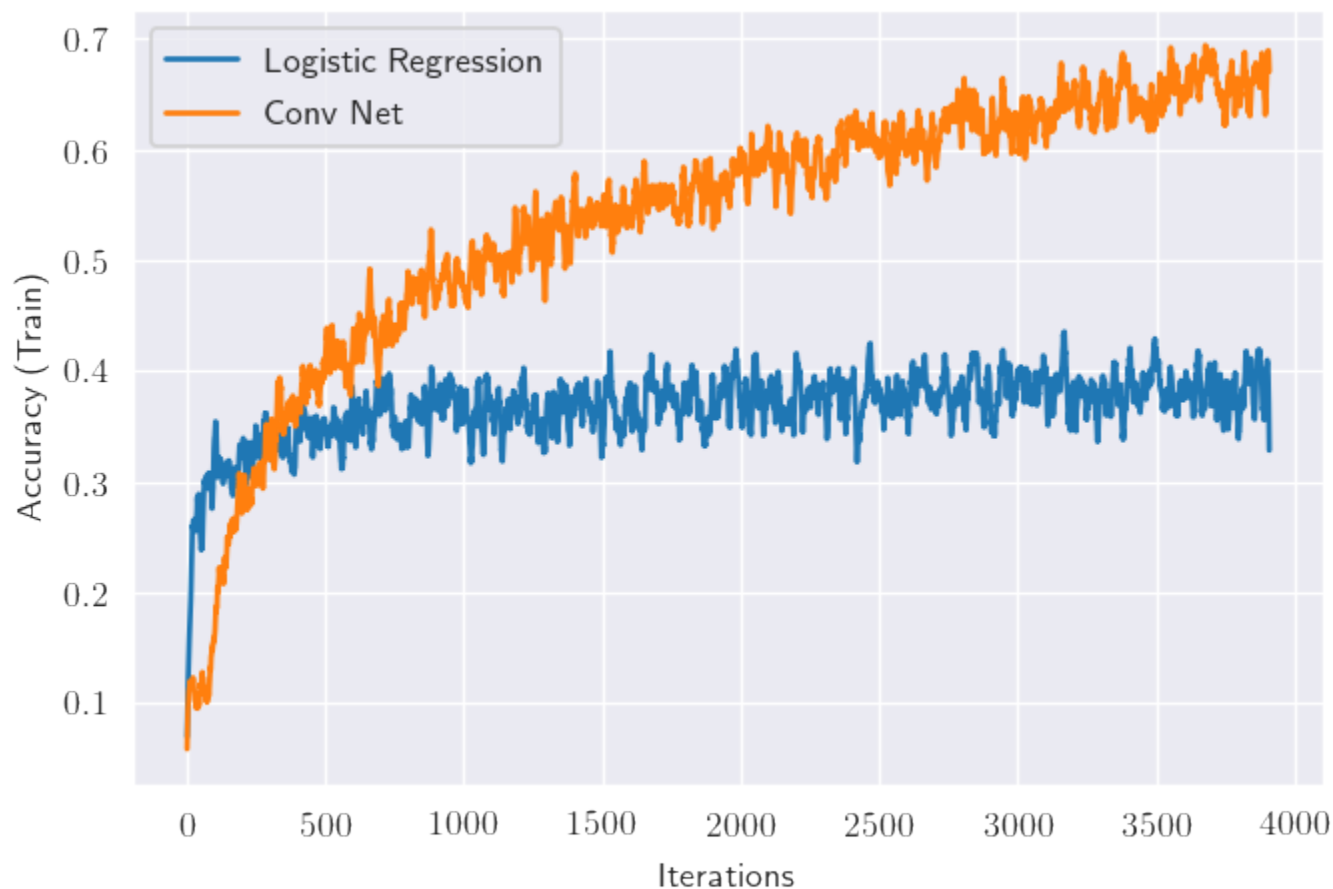
Number of Parameters




Question (Extension): Calculate the number of parameters in the logistic regression model and the above convnet.

The diagram above may be a useful guide. You may also want to look at the `model.parameters()` method for each model.

CNN vs Logistic Regression



Autoencoder



Autoencoder

- Learn encoder $f : \mathcal{X} \rightarrow \mathcal{F}$ from original to latent space.
- Learn decoder $g : \mathcal{F} \rightarrow \mathcal{X}$ from latent to original space.
- Unsupervised mapping. $g \circ f$ should be close to identity. Minimize quadratic loss over data:

$$\min_{f,g} \sum_k \|x_k - g \circ f(x_k)\|^2 \quad (1)$$

